

Matrix Multiplication: Homework

Clone the Strassen's project template from

https://github.com/albertocasagrande/AD_strassen_template

and solve the following exercises.

1. Generalize the implementation to deal with non-square matrices.

The solution can be found in the branch `rectangular` of this repository, in the folder [Strassen_alg](#).

2. Improve the implementation of the Strassen's algorithm by reducing the memory allocations and test the effects on the execution time.

The solution can be found in the function `strassen_matrix_multiplication_best` (and consequently in the function `strassen_aux_best`), contained in the file `strassen.c` in the folder [Strassen_alg](#).

I used only 2 matrices `SA` and `SB` for the S matrices, instead of allocating 10 matrices, while for the P matrices I used 4 matrices (`PA`, `PB`, `PC`, `PD`) instead of 7. I firstly computed P_2, P_4, P_5 and P_6 in `PA`, `PB`, `PC`, `PD` respectively, to be able to compute C_{11} , then I computed P_1 in `PD` (so replacing P_6) to be able to compute C_{12} , then I computed P_3 in `PA` (so replacing P_2) to be able to compute C_{21} , lastly I computed P_7 in `PB` (so replacing P_4) to be able to compute the last matrix C_{22} .

I compiled and run the code on Ulysses cluster in Sissa, for both square and rectangular matrices. The output was

- o for square matrices:

```
1 ./strassen_test.x
2
3      n  Naive Alg.  Strassen's Alg.  Str. Alg. best  Same result
4      1  0.000001    0.000020    0.000001      1 1
5      2  0.000000    0.000000    0.000000      1 1
6      4  0.000001    0.000000    0.000000      1 1
7      8  0.000009    0.000001    0.000001      1 1
8     16  0.000005    0.000004    0.000004      1 1
9     32  0.000032    0.000034    0.000028      1 1
10    64  0.000250    0.000244    0.000234      1 1
11   128  0.001934    0.001884    0.001853      1 1
12   256  0.100109    0.013079    0.012471      1 1
13   512  0.905262    0.089764    0.086891      1 1
14  1024  8.541473    0.627640    0.612996      1 1
15  2048  66.494477    4.420142    4.326053      1 1
16  4096 185.323830   31.036216   30.434407      1 1
```

- o for rectangular matrices:

```
1 ./strassen_test.x
2
3      dim  Naive Alg.  Strassen's Alg.  Str. Alg. best  Same result
```

4	1x	3x	4	0.000001	0.000022	0.000001	1 1
5	2x	6x	8	0.000001	0.000001	0.000000	1 1
6	4x	12x	16	0.000002	0.000001	0.000001	1 1
7	8x	24x	32	0.000007	0.000006	0.000006	1 1
8	16x	48x	64	0.000061	0.000049	0.000049	1 1
9	32x	96x	128	0.000390	0.000368	0.000364	1 1
10	64x	192x	256	0.022256	0.022456	0.022060	1 1
11	128x	384x	512	0.172694	0.174371	0.174951	1 1
12	256x	768x	1024	1.452749	0.465220	0.463335	1 1
13	512x	1536x	2048	12.976235	1.766585	1.750519	1 1
14	1024x	3072x	4096	69.019623	8.170306	8.084943	1 1

While on the new partition `frontend-beta` we have the following results:

- o for square matrices:


```
1 ./strassen_test.x
2
3      n  Naive Alg.  Strassen's Alg.  Str. Alg. best  Same result
4      1  0.000002   0.000006   0.000001   1 1
5      2  0.000001   0.000001   0.000001   1 1
6      4  0.000001   0.000001   0.000001   1 1
7      8  0.000003   0.000002   0.000002   1 1
8     16  0.000008   0.000006   0.000006   1 1
9     32  0.000047   0.000041   0.000041   1 1
10    64  0.000371   0.000347   0.000346   1 1
11   128  0.003007   0.002964   0.002960   1 1
12   256  0.025591   0.021378   0.020506   1 1
13   512  0.211942   0.109676   0.100745   1 1
14  1024  1.380608   0.741221   0.714642   1 1
15  2048 19.931641   5.106519   5.021268   1 1
16  4096 257.180293  35.810735  35.314176   1 1
```

- o for rectangular matrices:

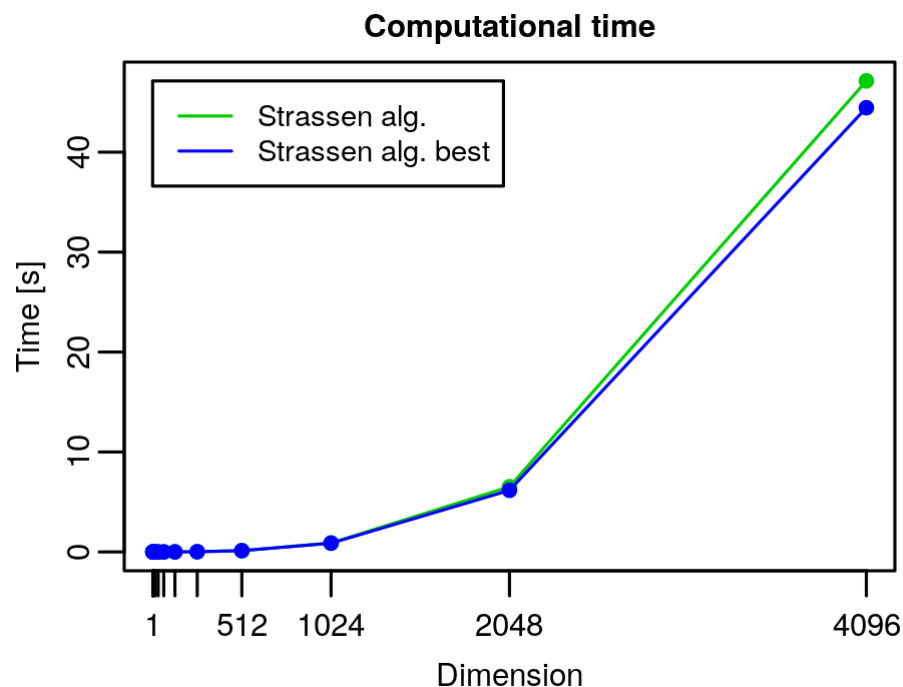
```
1 ./strassen_test.x
2
3      dim  Naive Alg.  Strassen's Alg.  Str. Alg. best  Same result
4     1x 3x  4  0.000002   0.000006   0.000001   1 1
5     2x 6x  8  0.000002   0.000001   0.000001   1 1
6     4x 12x 16 0.000003   0.000002   0.000002   1 1
7     8x 24x 32 0.000012   0.000009   0.000009   1 1
8    16x 48x 64 0.000072   0.000066   0.000066   1 1
9    32x 96x 128 0.000580   0.000552   0.000556   1 1
10   64x 192x 256 0.004737   0.004684   0.004678   1 1
11  128x 384x 512 0.039687   0.039087   0.039437   1 1
12  256x 768x 1024 0.294623   0.199306   0.196376   1 1
13  512x 1536x 2048 2.336765   1.403819   1.380727   1 1
14 1024x 3072x 4096 65.748678   9.800090   9.665999   1 1
```

We can see that the Strassen's algorithm is much better than the naive one, while the Strassen's algorithm with reduced memory allocations is only slightly better. Besides, it seems that on the new partition of Ulysses the times are a bit higher.

In the following graphs we can see that only in the last two/three points the time is significantly different from 0. Unfortunately, we have too few significant points to establish with certainty that the complexity of the naive algorithm is $\Theta(n^3) = \Theta(n^{\log_2 8})$ and the one of the Strassen's algorithm is $\Theta(n^{\log_2 7})$, even though the graph is growing very quickly. The problem is that with high power of 2 in n the matrices become very very big and are impossible to store in memory.



We can see that the Strassen's algorithm is much much more efficient than the naive algorithm.



Moreover, we can see that the Strassen's algorithm which uses only 6 matrices instead of 17 is also faster, beside being more memory efficient.